

OS

- Purpose: Hardware Abstraction; Resource Management (Time & Space Multiplexing)

User Mode VS Kernel Mode

- *User Mode*: cannot issue privileged instructions; can access parts of memory allowed by kernel mode code
- *Kernel Mode*: can issue all instructions; access all memory
- *Privileged*: instructions/memory locations that can interfere with other processes

I/O Devices

- Device → Controller → Driver
- Accessed via system calls (I/O calls are privileged instructions)
- Can be done using:
 - Busy Waiting: Driver sits in tight loop; poll device (see if done/busy)
 - Interrupts: Device sends interrupt to driver when done
 - Direct Memory Access (DMA): chip controls flow of bits between memory and controller without CPU intervention; chip causes interrupt when done

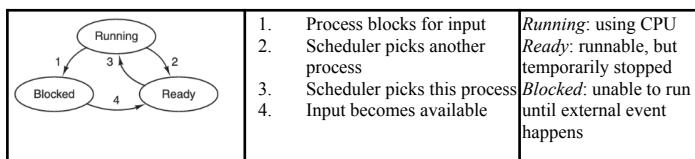
System Calls

- *System Call*: allows user mode program/process to request a service from OS kernel
 - switches user→kernel mode; significant performance overhead
 - Process, File, Directory Management; kill process, get time
- 1. Put the system call number in a place where OS expects it (register)
- 2. Execute a TRAP instruction to switch from user mode to kernel mode and start execution at a fixed address within the kernel
- 3. The kernel code that starts following the TRAP examines the system-call number and dispatches to the correct system-call handle
- 4. System call handler runs
- 5. Once it has completed its work, control may be returned to the user program
- 6. The system call may block the caller, then the OS will look around to see if some other process can be run next

Inter-Process Communication

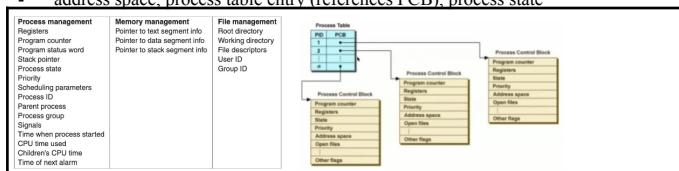
Process

- *Process*: abstraction of a running program; encapsulates all resources/information needed to run a program



Process Table

- linked list of PCBs (Process Control Blocks)
 - *PCB*: contains state required to resume a process
- Current process runs on RAM
- Not-currently-running process: info stored in process table
- address space, process table entry (references PCB), process state



Threads

- *Thread*: sequential execution stream within the process; entities scheduled for execution by CPU
- PRO: parallel; shared address space; reduced overhead (easier create, destroy, switch, communicate); □ I/O
- CON: not help performance when CPU bound (more threads than CPU cores)

Per Process Items (shared among threads)	Per Thread
<ul style="list-style-type: none"> - Address space - Global Variables - Open files - Child processes - Pending alarms - Signals and signal handlers - Accounting information 	<ul style="list-style-type: none"> - Program counter - Registers - Stack - State <p>Threads can read, write or delete other thread's stacks</p>

Multiprogramming

- allows multiple programs to share CPU and run in (pseudo)parallel
- but CPU can only run one program at a time – OS gives illusion of parallelism by switching between processes quickly
 - processes are sequential but execution is interleaved
 - programs cannot rely on timing assumptions

Race Conditions

- when multiple processes/threads read/write from common storage

<pre>stack //global shared variable void my_stack_function() { if (isEmpty(stack)) return; A: int s = pop(stack); //do something with s... } Another potential execution transcript 1. A tests isEmpty(stack) => false 2. B tests isEmpty(stack) => false 3. A pops => stack is now empty 4. B pops -> Exception?</pre>	<p>Avoiding Race Conditions</p> <ul style="list-style-type: none"> - <i>Mutual Exclusion</i>: ensure that 1 process exists in the critical region at any given time - <i>Critical Region (CR)</i>: part of program where shared resource/data is accessed - Conditions to achieve good solution <ol style="list-style-type: none"> 1. No two processes may be <u>simultaneously</u> inside their CRs 2. <u>No assumptions</u> can be made about speed and number of CPUs 3. No process running <u>outside</u> its CR can <u>block</u> other processes
--	--

	4. No process should have to wait forever to enter CR
--	---

Deadlocks

4 Conditions for deadlock

1. Mutual Exclusion: each resource can only be assigned to 1 process at a time
2. Hold-and-Wait: process currently holding resources can request new resources
3. No-Preemption: resources previously granted cannot be forcibly taken away
4. Circular Wait: circular list of 2 or more processes; each waiting for resource held by next member of chain

Dealing with Deadlocks

1. Ignore the problem
2. Detection (Resource allocation graph and graph algorithms) and recovery
3. Avoid by careful resource allocation
4. Prevention

Process Scheduling

- *Scheduling Algorithm*: determines which process to run next

- General Scheduling Goals
 - *Fairness*: giving each process a fair share of the CPU
 - *Policy Enforcement*: seeing that stated policy is carried out
 - *Balance*: keeping all parts of the system busy
- When to schedule: process created/exits/blocks, I/O interrupt, clock interrupt

Preemptive VS Non-preemptive Scheduling

- *Preemptive*: pick a process and let it run for a quantum (maximum fixed amount of time); clock interrupt at the end of time interval to give control to CPU back to scheduler
- *Non-preemptive*: pick process and let it run for as long as it wants until blocks/voluntarily releases to CPU
- Context Switching: happens every time OS switches process; expensive

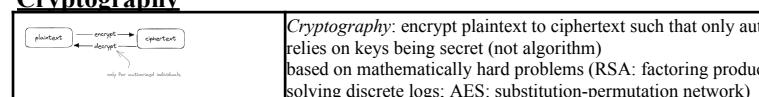
Scheduling Environments

Batch	Interactive	Real Time
-Don't require user interaction -A set of (potentially periodic) jobs that need to run	-Interactive users expect fast response times/responsiveness from the system	-Need to meet deadlines -Not really focused in this subject
Can be Preemptive or Non-Preemptive	Preemptive -by interrupting the process after a fixed time, it ensures that every process will produce a fast response -rather than having to wait, a very long time for one process to finish before moving onto the next one, which may lead to the user having to wait and thus not achieving fast interactivity	Can be Preemptive or Non-Preemptive -possibly schedule the process with fastest deadline first
	-Throughput: maximise jobs per hour -Turnaround time: minimise time between submission and termination -CPU utilisation: keep CPU busy all the time	-Meeting deadlines: avoid losing data -Predictability: avoid quality degradation in multimedia systems
-[Non-Pre] First-Come First Served (FCFS) -[Non-Pre] Shortest Job First (SJF) -[Pre] Shortest Remaining Time Next (SRT)	-[Pre] Round Robin	-[Pre] Priority Scheduling

Memory Management

- PRO: support multiprogramming; security (isolation between process and OS memory); enable running processes with large memory requirements
- *Memory Manager*: manages (part of) memory hierarchy; keep track of which parts of memory are in use; allocate when needed – deallocate when done
- *Swapping*: bringing in each process in its entirety, running it for a while

Cryptography

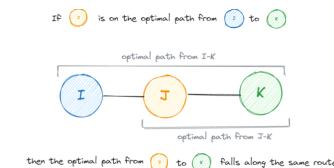


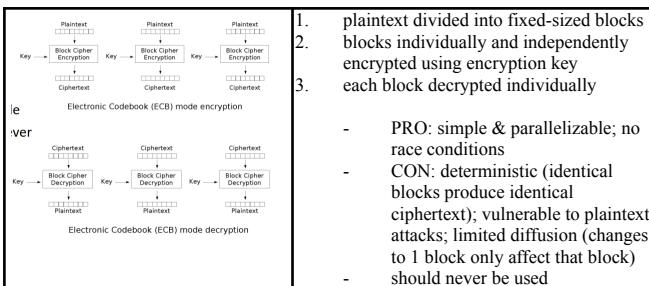
CIA

- *Confidentiality*: only sender and intended receiver should understand contents of transmitted message
- *Authentication*: establish identities of one or both endpoints
- *Integrity*: ensure that messages are not altered

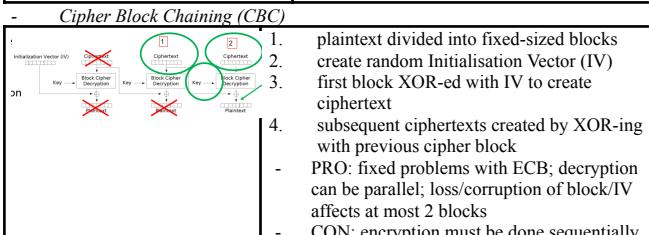
Symmetric Cryptography

- Same key used for encryption and decryption
- *Advanced Encryption Standard (AES)*
 - Break data into blocks and encrypt each block
 - Modes of operation: ECB, CBC – determines how each block is treated/linked
 - *Electronic Codebook (ECB)*





1. plaintext divided into fixed-sized blocks
 2. blocks individually and independently encrypted using encryption key
 3. each block decrypted individually
- PRO: simple & parallelizable; no race conditions
 - CON: deterministic (identical blocks produce identical ciphertext); vulnerable to plaintext attacks; limited diffusion (changes to 1 block only affect that block) should never be used

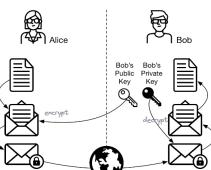


1. plaintext divided into fixed-sized blocks
 2. create random Initialization Vector (IV)
 3. first block XOR-ed with IV to create ciphertext
 4. subsequent ciphertexts created by XOR-ing with previous cipher block
- PRO: fixed problems with ECB; decryption can be parallel; loss/corruption of block/IV affects at most 2 blocks
 - CON: encryption must be done sequentially

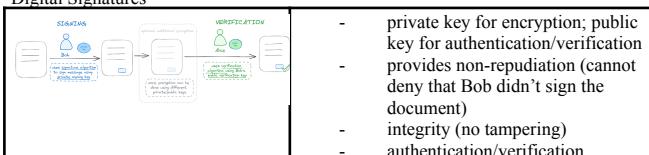
Symmetric Cryptography

- Two different keys for encryption and decryption
- CONS: slower, not suitable for large amounts of data/multiple blocks

- Alice generates her key pair (**PublicKey**, **PrivateKey**)
- She posts her **PublicKey** online
- Bob generates a **SK** with symmetric encryption
- Bob runs **CipherText := Encrypt(PublicKey, SK)**
- Bob sends the **CipherText** to Alice
- Alice runs **SK := Decrypt(PrivateKey, CipherText)**
- Now Alice and Bob share a **SK** and can exchange Symmetrically encrypted messages
- This is sort of how TLS (HTTPS) works
- A lot of further detail, but in principle, TLS is a key exchange protocol

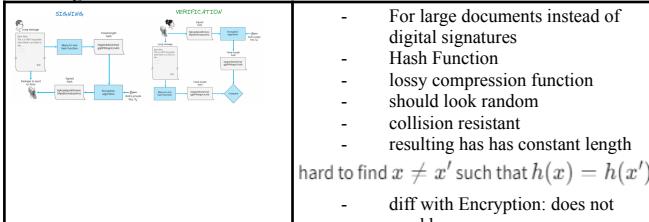


Digital Signatures



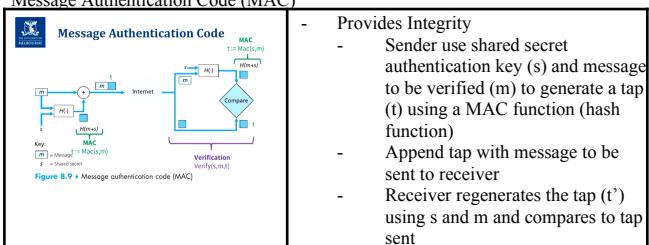
- private key for encryption; public key for authentication/verification
- provides non-repudiation (cannot deny that Bob didn't sign the document)
- integrity (no tampering)
- authentication/verification

Hash Digest



- For large documents instead of digital signatures
- Hash Function
- lossy compression function
- should look random
- collision resistant
- resulting has has constant length
- hard to find $x \neq x'$ such that $h(x) = h(x')$
- diff with Encryption: does not need key; one-way

Message Authentication Code (MAC)



- Provides Integrity
- Sender use shared secret authentication key (s) and message to be verified (m) to generate a tap (t) using a MAC function (hash function)
- Append tap with message to be sent to receiver
- Receiver regenerates the tap (t') using s and m and compares to tap sent

Authenticated Encryption – Integrity + Confidentiality



- General construction: Encrypt-then-Mac:
 - $m := Encrypt(SK, m)$
 - $t := Mac(SK, m)$
- Verify: If $Verify(s, t, c)$ returns 0, do not decrypt
- Examples: AES-GCM, AES-OCB, AES-CCM

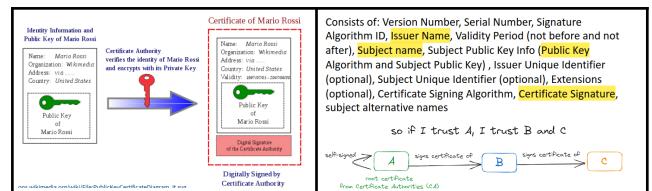
Diffie Hellman Key Exchange – Sharing Secret Keys



1. Alice and Bob agree upon p
2. Alice
 - a. picks random value s
 - b. computes $X = g^s \text{ mod } p$
 - c. sends X to Bob
3. Bob
 - a. picks random value t
 - b. computes $Y = g^t \text{ mod } p$
4. Alice calculates secret
 - $s = g^{ts} \text{ mod } p = g^t \text{ mod } p$
5. Bob calculates secret
 - $s = X^t \text{ mod } p = g^{ts} \text{ mod } p$
6. Both Alice and Bob have the same secret
 - $s = g^{ts} \text{ mod } p = g^{t^s} \text{ mod } p$

Public Key Infrastructure (Certificates) – Authentication

- Certificate: bind public key to identity of an entity (person, organisation, website etc.); signed by trusted entities (Certificate Authorities: explicitly trusted; predefined in OS; root cert is self-signed)



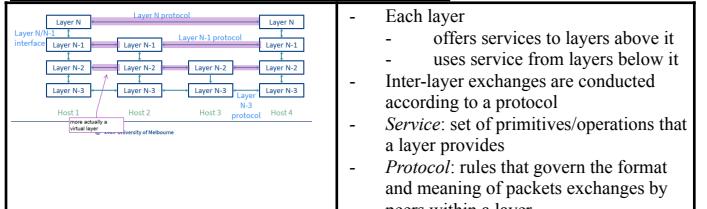
Consists of: Version Number, Serial Number, Signature, Algorithm ID, Issuer Name, Validity Period (not before and not after), Subject name, Subject Public Key Info [Public Key Algorithm and Subject Public Key], Issuer Unique Identifier (optional), Subject Unique Identifier (optional), Extensions (optional), Certificate Signing Algorithm, CertificateSignature, subject alternative name

so if I trust A, I trust B and C

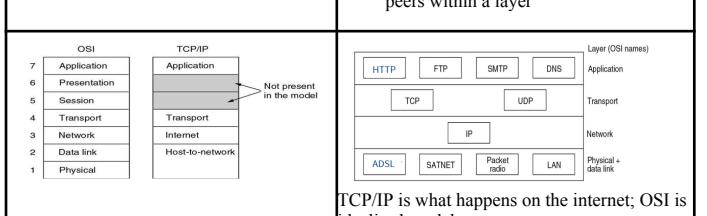
self-signed \xrightarrow{A} signs certificate of \xrightarrow{B} signs certificate of \xrightarrow{C}

root certificates from certificate authorities (CA)

Network Protocols and Service Models

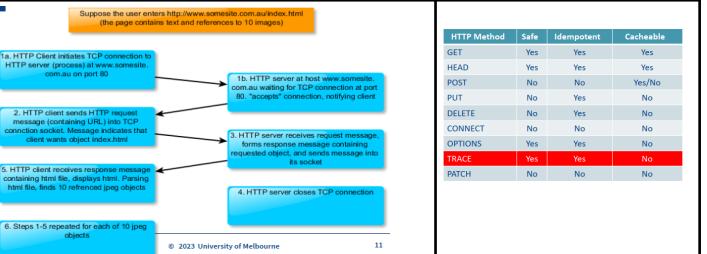


- Each layer
 - offers services to layers above it
 - uses service from layers below it
- Inter-layer exchanges are conducted according to a protocol
- **Service**: set of primitives/operations that a layer provides
- **Protocol**: rules that govern the format and meaning of packets exchanges by peers within a layer



TCP/IP is what happens on the internet; OSI is idealised model

HTTP



HTTP Method	Safe	Idempotent	Cacheable
GET	Yes	Yes	Yes
HEAD	Yes	Yes	Yes
POST	No	No	Yes/No
PUT	No	Yes	No
DELETE	No	Yes	No
CONNECT	No	No	No
OPTIONS	Yes	Yes	No
TRACE	Yes	Yes	No
PATCH	No	No	No

- Non-Persistent: requires 2 response times (initiate TCP; initial HTTP request) per object + file transmission time; OS overhead for each TCP connection
- Persistent: server leaves connection open after sending response
- Idempotent: multiple identical requests have same effect
- Safe: only for information retrieval, should not change state

Codes&Headers

Code	Meaning	Examples
1xx	Information	100 – server agrees to handle client's request
2xx	Success	200 = request succeeded; 204 = no content present
3xx	Redirection	301 = page moved; 304 = cached page still valid
4xx	Client error	403 = forbidden page; 404 = page not found
5xx	Server error	500 = internal server error; 503 try again later

Header	Type	Description
User-Agent	Request	Information about the browser and its platform
Accept	Request	The type of pages the client can handle
Accept-Charset	Request	The character sets that are acceptable to the client
Accept-Encoding	Request	The compression formats the client can handle
Accept-Language	Request	The natural languages the client can handle
If-Modified-Since	Request	Time and data to check freshness
If-None-Match	Request	Previously sent tags to check freshness
Host	Request	The server's DNS name
Authorization	Request	A list of the client's credentials
Referer	Request	The previous URL from which the request came
Cookie	Request	Previously sent cookie back to the server
Set-Cookie	Response	Cookie for the client to store
Server	Response	Information about the server

Header	Type	Description
Content-Encoding	Response	How the content is encoded (e.g., gzip)
Content-Language	Response	The natural language used in the page
Content-Length	Response	The page's length in bytes
Status-Line	(protocol status code and phrase)	HTTP/1.1 200 OK
Header-lines		Date: Thu, 01 Aug 2009 12:00:15 GMT
Header-lines		Server: Apache/2.2.11 (Ubuntu)
Header-lines		Last-modified: Mon, 22 Jun 2009
Header-lines		Content-Length: 6821
Header-lines		Content-Type: text/html
Data, e.g., requested HTML file		<html><head>
Header		</html>

Request & Response

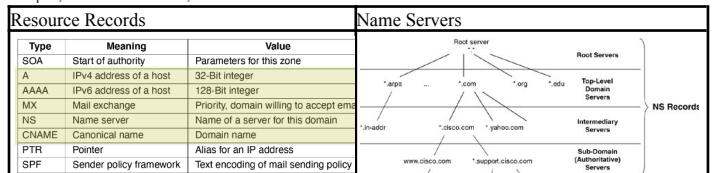
request line (get, post, head)	GET /onefile.php.html HTTP/1.1 Host: www.somefile.com.au User-Agent: Mozilla/4.0 Connection: close Accept-Language: fr
header lines	Date: Thu, 01 Aug 2009 12:00:15 GMT Server: Apache/2.2.11 (Ubuntu) Last-modified: Mon, 22 Jun 2009
status line (protocol status code and phrase)	HTTP/1.1 200 OK
header lines	Content-Type: text/html
data, e.g., requested HTML file	<html><head>

DNS

Domain Name System (4 elements): Domain name space, name servers, DNS database, resolvers

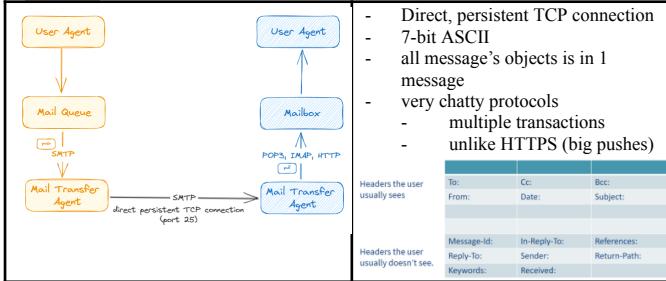
- You want to look up an address
- 1. You contact the resolver
- 2. Resolver will contact multiple name servers
- 3. The name servers will tell you something about the database
- 4. The structure of the database is captured by the DNS on the domain name space

Domain Name characteristics: not case sensitive, up to 63 chars per constituent, 255 chars per path, can be ASCII/UNICODE



- Root Name Server: 13 globally; authoritative cluster for enquiries
- Top-level Domain Servers: (com, edu, gov, int, mil, net, org, ...)
- Authoritative Servers: maintained by organisation/service provider; takes domain and subdomain → returns correct IP address
- Local DNS Server: each ISP/company/university; returns cached value if exists; acts as proxy, forwards requests up hierarchy

SMTP



Note: many back-and-forth exchanges.

S: 250 Hello crepes.fr, pleased to meet you

C: 250 Alice@crepes.fr... Sender ok

S: 250 bob@hamburger.edu ... Recipient ok

C: DATA

S: 354 Enter mail, end with "." on a line by itself

C: Do you like ketchup?

S: How about pickles?

Modern approach

C: 250 Message accepted for delivery

C: QUIT

S: 221 hamburger.edu closing connection

Multipurpose Internet Mail Extensions: Limited to English and not for other languages or audio/img

MIME has 6 additional message headers:

- MIME-version: identifies the MIME VERSION

- Content-Description: human readable describing contents

- Content-Id: unique identifier

- Content-Transfer-Encoding: how body is wrapped for transmission

- Content-Type: type and format of content

Type	Example subtypes	Description
text	plain, html, xml, css	Text in various formats
image	gif, jpeg, tiff	Pictures
audio	basic, mpeg, mp4	Sounds
video	mpeg, mp4, quicktime	Movies
model	vrml	3D model
application	octet-stream, pdf, javascript, zip	Data produced by applications
message	http://rfc22	Encapsulated message
multipart	mixed, alternative, parallel, digest	Combination of multiple types

Post Office Protocol(POP3 for delivery): hard to access email from multiple devices designed to keep mail until it's fetched, and then once it's fetched by a client, it's deleted from the server. That means you can't easily download it to the next device. The client then arranges it into folders, performs searching etc. If you configure it to use on two client computers, then putting it into a folder on one won't do the same on the other. Email cannot be reused.

Internet Message Access Protocol: IMAP keeps user state across sessions; Retain mailbox contents online (server) and allow manipulation of offline messages and mailbox folders; Implications of server infrastructure to support high volume of IMAP users. This implies storage projections by the provider, and hence limitations designed so that mail stays on the server, and the server has folders; The client can download email, and then delete it once it has been displayed (or cache it for offline viewing, of course, but the definitive copy is still on the server).

TCP

- Transport layer: tidy up end-to-end (lost/duplicated/corrupted packets)
- stream-oriented (segmenting into IP datagrams), reliable (dropped/duplicated) and in-order protocol
- **full duplex:** data in both directions, simultaneously
- **end-to-end:** exact pairs of senders and receivers
- **byte streams:** not message streams; message boundaries not preserved (no way of knowing sender has ended message)
- **buffer capable:** TCP entity can choose to buffer prior to sending or not; buffering reduces overhead (\square headers) but \square delay
- **segments:** 20-60 byte header + zero or more data bytes (TCP decides size of segments)
- **5-tuple address:** sender IP, receiver IP, sender port#, receiver port#, protocol

Primitives		
LISTEN	Packet Sent	Meaning
CONNECT	(none)	Block until something tries to connect
SEND	DATA	Send information
RECEIVE	(none)	Block until DATA packet arrives
DISCONNECT	DISCONNECTION REQ.	This sends what release the connection
Status	Simplified name	Description
CLOSED	Idle	No connection is active or pending
LISTEN	Pass. est.	The server is waiting for an incoming call
SYN/REQ	Pass. est.	A connection request has arrived; wait for ACK
SYN/SENT	Act. est.	The application has started to open a connection
ESTABLISHED	Established	The normal data transfer state
FIN/WAIT 1	Act. disc.	The application has said it is finished
FIN/WAIT 2	Act. disc.	The other side has agreed to release
TIME WAIT	Act. disc.	Wait for all packets to die off
CLOSING	Act. disc.	Both sides have tried to close simultaneously
CLOSE-WAIT	Pass. disc.	The other side has initiated a release
LAST ACK	Pass. disc.	Wait for all packets to die off

Sockets



Socket Finite State Machine

Bold before slash:
System call e.g., connect

Non-bold before slash:
Packet received e.g., SYN

After slash:
Packet sent e.g., SYN

Don't memorize this.
Just know it exists.

Connection request segment received

Connect primitive executed

ESTABLISHED

Connection accepted segment received

Disconnect primitive executed

ACTIVE DISCONNECT PENDING

Disconnection request segment received

Disconnected primitive executed

IDLE

Disconnection request segment received

Disconnected primitive executed

ACTIVE DISCONNECT PENDING

Disconnection request segment received

Disconnected primitive executed

IDLE

Connection request segment received

Connect primitive executed

ESTABLISHED

Connection accepted segment received

Disconnect primitive executed

ACTIVE DISCONNECT PENDING

Disconnection request segment received

Disconnected primitive executed

IDLE

Connection request segment received

Connect primitive executed

ESTABLISHED

Connection accepted segment received

Disconnect primitive executed

ACTIVE DISCONNECT PENDING

Disconnection request segment received

Disconnected primitive executed

IDLE

Connection request segment received

Connect primitive executed

ESTABLISHED

Connection accepted segment received

Disconnect primitive executed

ACTIVE DISCONNECT PENDING

Disconnection request segment received

Disconnected primitive executed

IDLE

Connection request segment received

Connect primitive executed

ESTABLISHED

Connection accepted segment received

Disconnect primitive executed

ACTIVE DISCONNECT PENDING

Disconnection request segment received

Disconnected primitive executed

IDLE

Connection request segment received

Connect primitive executed

ESTABLISHED

Connection accepted segment received

Disconnect primitive executed

ACTIVE DISCONNECT PENDING

Disconnection request segment received

Disconnected primitive executed

IDLE

Connection request segment received

Connect primitive executed

ESTABLISHED

Connection accepted segment received

Disconnect primitive executed

ACTIVE DISCONNECT PENDING

Disconnection request segment received

Disconnected primitive executed

IDLE

Connection request segment received

Connect primitive executed

ESTABLISHED

Connection accepted segment received

Disconnect primitive executed

ACTIVE DISCONNECT PENDING

Disconnection request segment received

Disconnected primitive executed

IDLE

Connection request segment received

Connect primitive executed

ESTABLISHED

Connection accepted segment received

Disconnect primitive executed

ACTIVE DISCONNECT PENDING

Disconnection request segment received

Disconnected primitive executed

IDLE

Connection request segment received

Connect primitive executed

ESTABLISHED

Connection accepted segment received

Disconnect primitive executed

ACTIVE DISCONNECT PENDING

Disconnection request segment received

Disconnected primitive executed

IDLE

Connection request segment received

Connect primitive executed

ESTABLISHED

Connection accepted segment received

Disconnect primitive executed

ACTIVE DISCONNECT PENDING

Disconnection request segment received

Disconnected primitive executed

IDLE

Connection request segment received

Connect primitive executed

ESTABLISHED

Connection accepted segment received

Disconnect primitive executed

ACTIVE DISCONNECT PENDING

Disconnection request segment received

Disconnected primitive executed

IDLE

Connection request segment received

Connect primitive executed

ESTABLISHED

Connection accepted segment received

Disconnect primitive executed

ACTIVE DISCONNECT PENDING

Disconnection request segment received

Disconnected primitive executed

IDLE

Connection request segment received

Connect primitive executed

ESTABLISHED

Connection accepted segment received

Disconnect primitive executed

ACTIVE DISCONNECT PENDING

Disconnection request segment received

Disconnected primitive executed

IDLE

Connection request segment received

Connect primitive executed

ESTABLISHED

Connection accepted segment received

Disconnect primitive executed

ACTIVE DISCONNECT PENDING

Disconnection request segment received

Disconnected primitive executed

IDLE

Connection request segment received

Connect primitive executed

ESTABLISHED

Connection accepted segment received

Disconnect primitive executed

ACTIVE DISCONNECT PENDING

Disconnection request segment received

Disconnected primitive executed

IDLE

Connection request segment received

Connect primitive executed

ESTABLISHED

Connection accepted segment received

Disconnect primitive executed

ACTIVE DISCONNECT PENDING

Disconnection request segment received

Disconnected primitive executed

IDLE

Connection request segment received

Connect primitive executed

ESTABLISHED

Connection accepted segment received

Disconnect primitive executed

ACTIVE DISCONNECT PENDING

Disconnection request segment received

Disconnected primitive executed

IDLE

Connection request segment received

Connect primitive executed

ESTABLISHED

Connection accepted segment received

Disconnect primitive executed

ACTIVE DISCONNECT PENDING

Disconnection request segment received

Disconnected primitive executed

IDLE

Connection request segment received

Connect primitive executed

ESTABLISHED

Connection accepted segment received

Disconnect primitive executed

ACTIVE DISCONNECT PENDING

Disconnection request segment received

Disconnected primitive executed

IDLE

Connection request segment received

Connect primitive executed

ESTABLISHED

Connection accepted segment received

Disconnect primitive executed

ACTIVE DISCONNECT PENDING

Disconnection request segment received

Disconnected primitive executed

IDLE

Connection request segment received

Connect primitive executed

ESTABLISHED

Connection accepted segment received

Disconnect